

# Tarefa#2 - Aprendizado supervisionado - MO432

June 8, 2021

Universidade Estadual de Campinas (UNICAMP), Instituto de Computação (IC)

Prof. Jacques Wainer, 2021s1

```
[1]: # RA & Name
print('265673: ' + 'Gabriel Luciano Gomes')
print('264965: ' + 'Décio Luiz Gazzoni Filho')
print('192880: ' + 'Lucas Borges Rondon')
```

```
265673: Gabriel Luciano Gomes
264965: Décio Luiz Gazzoni Filho
192880: Lucas Borges Rondon
```

## 1 Leitura da base de dados

```
[2]: import pandas as pd
import requests
import io

url = "https://archive.ics.uci.edu/ml/machine-learning-databases/00514/
↳Bias_correction_ucl.csv"
s = requests.get(url).content

db = pd.read_csv(io.StringIO(s.decode('utf-8')))
db.head()
```

```
[2]:      station      Date  Present_Tmax  ...  Solar radiation  Next_Tmax  Next_Tmin
0         1.0  2013-06-30          28.7  ...      5992.895996         29.1         21.2
1         2.0  2013-06-30          31.9  ...      5869.312500         30.5         22.5
2         3.0  2013-06-30          31.6  ...      5863.555664         31.1         23.9
3         4.0  2013-06-30          32.0  ...      5856.964844         31.7         24.3
4         5.0  2013-06-30          31.4  ...      5859.552246         31.2         22.5
```

[5 rows x 25 columns]

```
[3]: print(f'A base de dados possui {db.shape[0]} instâncias com {db.shape[1]}
↳atributos.')
```

A base de dados possui 7752 instâncias com 25 atributos.

## 2 Pré-processamento dos dados

### 2.1 Remoção das colunas “Next\_Tmin” e “Date”

```
[4]: db = db.drop(columns=['Next_Tmin', 'Date'])  
  
db.head()
```

```
[4]:
```

	station	Present_Tmax	Present_Tmin	...	Slope	Solar radiation	Next_Tmax
0	1.0	28.7	21.4	...	2.7850	5992.895996	29.1
1	2.0	31.9	21.6	...	0.5141	5869.312500	30.5
2	3.0	31.6	23.3	...	0.2661	5863.555664	31.1
3	4.0	32.0	23.4	...	2.5348	5856.964844	31.7
4	5.0	31.4	21.9	...	0.5055	5859.552246	31.2

[5 rows x 23 columns]

### 2.2 Identificação de valores nulos

```
[5]: db.isna().sum() # Verificar se existem valores desconhecidos na base de dados
```

```
[5]: station                2  
Present_Tmax              70  
Present_Tmin              70  
LDAPS_RHmin               75  
LDAPS_RHmax               75  
LDAPS_Tmax_lapse          75  
LDAPS_Tmin_lapse          75  
LDAPS_WS                  75  
LDAPS_LH                   75  
LDAPS_CC1                  75  
LDAPS_CC2                  75  
LDAPS_CC3                  75  
LDAPS_CC4                  75  
LDAPS_PPT1                 75  
LDAPS_PPT2                 75  
LDAPS_PPT3                 75  
LDAPS_PPT4                 75  
lat                        0  
lon                        0  
DEM                        0  
Slope                      0  
Solar radiation            0  
Next_Tmax                  27  
dtype: int64
```

### 2.2.1 Remoção dos valores nulos

```
[6]: db = db.dropna()
      db.shape
```

```
[6]: (7588, 23)
```

### 2.2.2 Separação da base de dados (entrada e saída)

```
[7]: X = db.drop(columns=['Next_Tmax'])
      y = db['Next_Tmax']
```

### 2.2.3 Centering e Scaling dos dados de entrada

```
[8]: from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler()
      X = scaler.fit_transform(X)
```

## 3 Processamento do conjunto de dados

### 3.1 Imports necessários

```
[9]: from random import sample
      from scipy.stats import loguniform, uniform
      from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor
      from sklearn.exceptions import ConvergenceWarning
      from sklearn.linear_model import Lasso, LinearRegression, Ridge
      from sklearn.model_selection import cross_val_score, GridSearchCV,
      ↪RandomizedSearchCV
      from sklearn.neighbors import KNeighborsRegressor
      from sklearn.neural_network import MLPRegressor
      from sklearn.svm import SVR
      from sklearn.tree import DecisionTreeRegressor
      from statistics import mean
```

```
[10]: def fit_values(model, X, y):
        cvs = cross_val_score(model,
                               X,
                               y,
                               cv=5,
                               scoring='neg_root_mean_squared_error')

        return cvs

      def find_hyperparams(model, space, X, y, n_iter):
        search = RandomizedSearchCV(model,
                                     space,
```

```

        n_iter = n_iter,
        scoring = 'neg_root_mean_squared_error',
        n_jobs = -1,
        cv = 5,
        random_state = 1)

    return search.fit(X,y)

def best_results(model, space, X, y, hparams_name, n_iter = 10):
    hparams = find_hyperparams(model, space, X, y, n_iter)

    best_hparams = hparams.best_params_
    best_rmse     = -hparams.best_score_

    string_result = f'Melhor RMSE: {best_rmse:5f} para'

    for x in hparams_name:
        string_result += f' {x}: {best_hparams[x]:5f} \t'

    print(string_result)

    model_default = fit_values(model, X, y)

    print(f'RMSE para modelo default: {mean(-model_default):5f}')

    return (best_hparams, best_rmse, model_default)

```

### 3.2 Modelo Linear

```

[11]: linear_ = fit_values(LinearRegression(), X, y)

print('*' *10 + ' Modelo Linear ' + '*' *10)
print(f'Melhor RMSE: {mean(-linear_)}')

```

```

***** Modelo Linear *****
Melhor RMSE: 1.5775462124225403

```

### 3.3 Modelo Linear com regularização L2

```

[12]: space = dict()
space['alpha'] = loguniform(10**-3, 10**3)

print('*' *10 + ' Modelo Ridge ' + '*' *10)
ridge_ = best_results(Ridge(),
                      space,
                      X,

```

```
y,  
['alpha'])
```

\*\*\*\*\* Modelo Ridge \*\*\*\*\*

Melhor RMSE: 1.576248 para alpha: 20.986836

RMSE para modelo default: 1.577477

### 3.4 Modelo Linear com regularização L1

```
[13]: space = dict()  
space['alpha'] = loguniform(10**-3, 10**3)  
  
print('*' *10 + ' Modelo Lasso ' + '*' *10)  
lasso_ = best_results(Lasso(),  
                      space,  
                      X,  
                      y,  
                      ['alpha'])
```

\*\*\*\*\* Modelo Lasso \*\*\*\*\*

Melhor RMSE: 1.570089 para alpha: 0.013109

RMSE para modelo default: 2.024356

### 3.5 Modelo SVM Linear

```
[14]: space = dict()  
space['epsilon'] = [0.1, 0.3]  
space['C'] = loguniform(2**-5, 2**15)  
  
print('*' *10 + ' Modelo SVM Linear ' + '*' *10)  
svm_linear_ = best_results(SVR(kernel = 'linear', max_iter = 3000),  
                          space,  
                          X,  
                          y,  
                          ['epsilon', 'C'])
```

\*\*\*\*\* Modelo SVM Linear \*\*\*\*\*

/usr/local/lib/python3.7/dist-packages/sklearn/svm/\_base.py:231:

ConvergenceWarning: Solver terminated early (max\_iter=3000). Consider pre-processing your data with StandardScaler or MinMaxScaler.

% self.max\_iter, ConvergenceWarning)

Melhor RMSE: 1.788912 para epsilon: 0.300000 C: 0.824608

/usr/local/lib/python3.7/dist-packages/sklearn/svm/\_base.py:231:

ConvergenceWarning: Solver terminated early (max\_iter=3000). Consider pre-processing your data with StandardScaler or MinMaxScaler.

% self.max\_iter, ConvergenceWarning)

```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)

RMSE para modelo default: 1.843715

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)

```

### 3.6 Modelo SVM kernel RBF

```

[15]: space = dict()
space['epsilon'] = [0.1, 0.3]
space['C'] = loguniform(2**-5, 2**15)
space['gamma'] = loguniform(2**-9, 2**3)

print('*' *10 + ' Modelo SVM Linear ' + '*' *10)
svm_rbf_ = best_results(SVR(kernel = 'rbf', max_iter = 3000),
                        space,
                        X,
                        y,
                        ['epsilon', 'C', 'gamma'])

```

\*\*\*\*\* Modelo SVM Linear \*\*\*\*\*

```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)

```

Melhor RMSE: 1.566807 para epsilon: 0.100000 C: 14.611758 gamma: 0.002453

```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:231:
ConvergenceWarning: Solver terminated early (max_iter=3000). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
    % self.max_iter, ConvergenceWarning)

RMSE para modelo default: 1.685369

```

### 3.7 KNN

```

[16]: space = dict()
space['n_neighbors'] = list(range(1, 1001))

print('*' *10 + ' Modelo KNN ' + '*' *10)
svm_rbf_ = best_results(KNeighborsRegressor(),
                        space,
                        X,
                        y,
                        ['n_neighbors'])

```

```

***** Modelo KNN *****
Melhor RMSE: 1.845414 para n_neighbors: 38.000000
RMSE para modelo default: 1.927051

```

### 3.8 MLP

```

[17]: space = dict()
space['hidden_layer_sizes'] = list(range(5, 21, 3))

print('*' *10 + ' Modelo MLP ' + '*' *10)
mlp_ = best_results(MLPRegressor(),
                   space,
                   X,
                   y,
                   ['hidden_layer_sizes'],
                   6)

```

```

***** Modelo MLP *****

/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and

```

```

the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)

Melhor RMSE: 2.383525 para hidden_layer_sizes: 11.000000

/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)

RMSE para modelo default: 1.931995

/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)

```

```
[18]: mlp_
```

```

[18]: ({'hidden_layer_sizes': 11},
      2.3835245841742254,
      array([-2.00656122, -1.79094716, -2.02333629, -1.97491476, -1.8642152 ]))

```

### 3.9 Árvore de Decisão

```

[19]: space = dict()
      space['ccp_alpha'] = uniform(0.0, 0.04)

      print('*' *10 + ' Modelo de Árvore de Decisão ' + '*' *10)

      model = DecisionTreeRegressor()
      model.cost_complexity_pruning_path(X, y)

```



```
decision_tree_ = best_results(model,
                              space,
                              X,
                              y,
                              ['ccp_alpha'])
```

\*\*\*\*\* Modelo de Árvore de Decisão \*\*\*\*\*  
 Melhor RMSE: 1.861153 para ccp\_alpha: 0.021553  
 RMSE para modelo default: 2.208033

### 3.10 Random Forest

```
[20]: space = dict()
space['n_estimators'] = [10, 100, 1000]
space['max_features'] = [5, 10, 22]

print('*' *10 + ' Modelo Random Forest ' + '*' *10)

random_forest_ = best_results(RandomForestRegressor(),
                              space,
                              X,
                              y,
                              ['n_estimators', 'max_features'],
                              9)
```

\*\*\*\*\* Modelo Random Forest \*\*\*\*\*

/usr/local/lib/python3.7/dist-packages/joblib/externals/loky/process\_executor.py:691: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.  
 "timeout or by a memory leak.", UserWarning

Melhor RMSE: 1.627018 para n\_estimators: 1000.000000      max\_features: 10.000000  
 RMSE para modelo default: 1.658443

### 3.11 GBM

```
[21]: space = dict()
space['n_estimators'] = list(range(5, 101))
space['learning_rate'] = uniform(0.01, 0.3)
space['max_depth'] = [2,3]

print('*' *10 + ' Modelo GBM ' + '*' *10)

gbm_ = best_results(GradientBoostingRegressor(),
                   space,
                   X,
                   y,
```

```
['n_estimators', 'learning_rate', 'max_depth'])
```

\*\*\*\*\* Modelo GBM \*\*\*\*\*

Melhor RMSE: 1.585280 para n\_estimators: 77.000000 learning\_rate: 0.135107

max\_depth: 2.000000

RMSE para modelo default: 1.596553

### 3.12 Tabela Comparativa

Modelo	RMSE	Default	Melhores Hyperparâmetros
Linear	1.577546	-	-
Linear com L1	2.024356	1.570089	
Linear com L2	1.577477	1.576248	
SVM Linear	1.843715	1.788912	
SVM com brf	1.685369	1.566807	
KNN	1.927051	1.845414	
MLP	1.931995	2.383525	
Árvore de Decisão	2.208033	1.861153	
Random Forest	1.658443	1.627018	
GBM	1.596553	1.585280	